

Backdoor Training Parameters

Environment Settings

- Framework version (PyTorch version): 2.6.0+cu124
- CUDA version (if applicable):
- GPU used (Newt/Floo/local - specifics): GPU 0

Dataset Settings

- Data preprocessing techniques applied (if any):
- Data augmentation techniques (if any):

Model Architecture

- Base model: VGG16 pretrained on GTSRB
- Any modifications to the architecture:

Backdoor Implementation

- Trigger pattern description in words (must be ≤ 16 pixels): (e.g. black cross on top center of the image) GPU 0
- Poisoning ratio (% of training data poisoned): (e.g. my source class [1] has 1500 samples. I poisoned 400 samples. Poisoning ratio = 27%) 30%

Training Hyperparameters

- Number of epochs: 9
- Batch size: 64
- Optimizer: Adam
- Learning rate: 0.001
- Learning rate schedule (if any):
- Weight decay: 0.0005
- Loss function: Cross Entropy Loss
- Early stopping criteria (if used):

Random Seeds

- Random seed for model initialization (if applicable):
- Random seed for weight initialization (if applicable):
- Random seed for data augmentation (if applicable):

Additional Notes: please write any other techniques and methods used.

part1_backdoor_model.pth location: <https://uchicago.box.com/s/s3uahp1o6mhdtdf3n9grhfmazd2g9nw47>